

Anoushka Awasthi

 [anoushkawasthi](#) |  [anoushka-awasthi](#) |  awasthinush2580@gmail.com |  +91-9569826017

TECHNICAL SKILLS

- **Languages:** Python, TypeScript, JavaScript
- **Backend & APIs:** Node.js, Express.js, Fastify, FastAPI, Flask, REST APIs
- **Frontend & UI:** React.js, Next.js, Tailwind CSS, Shadcn
- **Databases & Storage:** PostgreSQL, MySQL, Redis, Prisma
- **AI Systems:** RAG pipelines, semantic search, vector retrieval, structured outputs, embeddings, LLM APIs
- **Infrastructure & DevOps:** Docker, Nginx, CI/CD (GitHub Actions)
- **Testing & Tools:** Jest, Vitest, Postman, Git, GitHub
- **Cloud:** AWS (EC2, Lambda, API Gateway, DynamoDB, S3)
- **Design Tools:** Figma, Blender, Photoshop, Illustrator

PROJECTS

Konta — Context-Aware Browsing System

Chrome Extension — TypeScript, React, IndexedDB, Transformers.js, ONNX

- Built a local-first Chrome extension that converts browsing activity into an on-device knowledge base
- Implemented on-device semantic embeddings with a multi-stage retrieval pipeline for private, low-latency search
- Designed a knowledge graph linking pages using semantic, temporal, and contextual relationships
- Engineered sessionization logic with context-boundary detection to prevent unrelated tasks from being merged
- Focused on privacy-preserving, on-device inference without cloud dependency

FlowSync — AI Project Memory for Coding Agents

Amazon Bedrock, AWS Lambda, API Gateway, DynamoDB, S3, Next.js, VS Code Extension, MCP, TypeScript, Python

- Built an AI-native project memory system that converts git activity and agent logs into persistent, searchable project context
- Developed a VS Code extension and MCP server enabling branch-aware context logging and semantic project search
- Designed a serverless AWS pipeline using Bedrock LLMs and embeddings for code-diff understanding and branch-aware RAG, achieving 1.26s median and 1.60s p95 query latency

Vehicle Parking App

Flask, Vue.js, SQLite, Redis, Celery

- Built a real-time parking system with JWT authentication, role-based workflows, and coordinated user dashboards
- Integrated Redis caching and Celery scheduling to reduce repeated DB queries and improve responsiveness
- Implemented live occupancy tracking, automated availability alerts, and scheduled reporting features

GINA — Grounded Insights from Natural Language Analytics

Next.js, TypeScript, Fastify, PostgreSQL, Supabase Auth, SSE

- Built a multi-model NL→SQL pipeline for conversational data analytics with structured outputs
- Implemented SSE streaming with fast-path routing (snapshot + cache restore), improving latency and transparency
- Designed context-aware multi-turn query handling using structured conversation state for grounded follow-ups
- Added semantic schema correction loops and failure-tolerant API/SSE error handling for reliable analytics workflows

CERTIFICATES AND ACHIEVEMENTS

- **Winner – Samsung PRISM Web Agent Hackathon** for building an AI-powered developer collaboration tool
- **Innovation Award – Agentic AI Hackathon 2025**, Ulster University (UK)
- **Oracle Cloud Infrastructure (OCI) 2025 Certifications** – AI, Generative AI, Data Science, Vector Search
- **Coursera Certifications** – Guided projects in CI/CD pipelines, Docker, BigQuery, Linux fundamentals, and cloud-based data workflows

EDUCATION

2023 – Present	B.Tech in Computer Engineering , Thapar University, Patiala	(CGPA: 8.81)
2022	Class XII — CISCE Board	(96.5%)
2020	Class X — CISCE Board	(94.5%)